

Class 07: Machine Learning 1


Kai Zhao(PID: A17599942)

Clustering

we will start today's lab with clustering methods, in particular so-called K-means. The main function for this in R is `kmeans()`

Let's try it on some made up data where we know that what the answer should be.

```
x <- rnorm(10000, mean=3)
hist(x)
```



60 points


```
tmp <- c(rnorm(30, mean=3), rnorm(30,mean=-3))
x <- cbind(x=tmp, y=rev(tmp))
x
```

| | x | y |
|-------|-----------|-----------|
| [1,] | 2.789181 | -3.246384 |
| [2,] | 1.785632 | -3.664413 |
| [3,] | 4.990367 | -2.262505 |
| [4,] | 4.004477 | -2.675746 |
| [5,] | 2.220653 | -2.385723 |
| [6,] | 2.820759 | -3.189301 |
| [7,] | 1.768562 | -1.531468 |
| [8,] | 3.499662 | -4.836351 |
| [9,] | 4.577516 | -3.541133 |
| [10,] | 3.276770 | -1.883271 |
| [11,] | 2.615736 | -2.793107 |
| [12,] | 2.183321 | -2.122602 |
| [13,] | 5.818104 | -1.512527 |
| [14,] | 4.225951 | -4.592156 |
| [15,] | 2.746062 | -3.534720 |
| [16,] | 3.781665 | -3.148096 |
| [17,] | 3.408777 | -2.320842 |
| [18,] | 3.624595 | -2.746581 |
| [19,] | 1.130785 | -2.267443 |
| [20,] | 3.505646 | -4.329642 |
| [21,] | 1.854807 | -3.827730 |
| [22,] | 4.408753 | -1.158637 |
| [23,] | 3.940465 | -2.591902 |
| [24,] | 4.282357 | -2.524787 |
| [25,] | 4.490019 | -3.639081 |
| [26,] | 2.712705 | -3.960268 |
| [27,] | 3.794480 | -3.933003 |
| [28,] | 1.588639 | -4.498693 |
| [29,] | 2.555489 | -5.383949 |
| [30,] | 3.158643 | -3.852315 |
| [31,] | -3.852315 | 3.158643 |
| [32,] | -5.383949 | 2.555489 |
| [33,] | -4.498693 | 1.588639 |
| [34,] | -3.933003 | 3.794480 |
| [35,] | -3.960268 | 2.712705 |
| [36,] | -3.639081 | 4.490019 |
| [37,] | -2.524787 | 4.282357 |

```
[38,] -2.591902  3.940465
[39,] -1.158637  4.408753
[40,] -3.827730  1.854807
[41,] -4.329642  3.505646
[42,] -2.267443  1.130785
[43,] -2.746581  3.624595
[44,] -2.320842  3.408777
[45,] -3.148096  3.781665
[46,] -3.534720  2.746062
[47,] -4.592156  4.225951
[48,] -1.512527  5.818104
[49,] -2.122602  2.183321
[50,] -2.793107  2.615736
[51,] -1.883271  3.276770
[52,] -3.541133  4.577516
[53,] -4.836351  3.499662
[54,] -1.531468  1.768562
[55,] -3.189301  2.820759
[56,] -2.385723  2.220653
[57,] -2.675746  4.004477
[58,] -2.262505  4.990367
[59,] -3.664413  1.785632
[60,] -3.246384  2.789181
```

We can pass this to the base R `plot()` function for a quick

```
plot(x)
```



```
k <- kmeans(x, centers=2, nstart=20)
k
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

| | x | y |
|---|-----------|-----------|
| 1 | 3.252019 | -3.131813 |
| 2 | -3.131813 | 3.252019 |

Clustering vector:

[illegible]

Within cluster sum of squares by cluster:

```
[1] 67.90501 67.90501
(between_SS / total_SS = 90.0 %)
```

Available components:

```
[1] "cluster"      "centers"      "tottss"       "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

Q1. How many points are in each cluster

k\$size

[1] 30 30

Q2. Cluster membership?

```
k$cluster
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```


Q3. Cluster centers?

k\$centers

| | x | y |
|---|-----------|-----------|
| 1 | 3.252019 | -3.131813 |
| 2 | -3.131813 | 3.252019 |


Q4 Plot my clustering results

```
plot(x, col= k$cluster, pch=16)
```



Q5 Cluster the data again with `kmean()` into 4 groups and plot the results.

```
k4 <- kmeans(x, center=4, nstart=20)
plot(x, col=k4$cluster, pch=16)
```



K-means is very popular mostly because it is fast and relatively straightforward to run and understand. It has a big limitation in that you need to tell it how many groups(k, or centers) you want.

hierarchical clustering

The main function in base R is called `hclust()`. You have to pass it in a “distance matrix” not just your input data.

You can generate a distance matrix with the `dist()` function.


```
hc <- hclust(dist(x))
hc
```

Call:

```
hclust(d = dist(x))
```

```
Cluster method : complete
Distance       : euclidean
Number of objects: 60
```


```
plot(hc)
```



To find the clusters (cluster membership vector) from a `hclust()` result we can “cut” the tree at a certain height that we like. For this we use the `cutree()` function.

```
plot(hc)  
abline(h=8,col="red")
```


Cluster Dendrogram




```
grps <- cutree(hc, h=8)
```

```
table(grps)
```

```
grps  
 1  2  
30 30
```

Q6 Plot our hclust results.

```
plot(x, col=grps, pch=16)
```



Principal Component Analysis

PCA of UK food data

Read data showing the consumption in grams(per person, per week) of 17 different types of food-stuff measured and averaged in the four countries of the United Kingdom.

Let's see how PCA can help us but first we can try conventional analysis.

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
x
```

| | | X | England | Wales | Scotland | N.Ireland |
|---|----------------|---|---------|-------|----------|-----------|
| 1 | Cheese | | 105 | 103 | 103 | 66 |
| 2 | Carcass_meat | | 245 | 227 | 242 | 267 |
| 3 | Other_meat | | 685 | 803 | 750 | 586 |
| 4 | Fish | | 147 | 160 | 122 | 93 |
| 5 | Fats_and_oils | | 193 | 235 | 184 | 209 |
| 6 | Sugars | | 156 | 175 | 147 | 139 |
| 7 | Fresh_potatoes | | 720 | 874 | 566 | 1033 |

| | | | | | |
|----|--------------------|------|------|------|------|
| 8 | Fresh_Veg | 253 | 265 | 171 | 143 |
| 9 | Other_Veg | 488 | 570 | 418 | 355 |
| 10 | Processed_potatoes | 198 | 203 | 220 | 187 |
| 11 | Processed_Veg | 360 | 365 | 337 | 334 |
| 12 | Fresh_fruit | 1102 | 1137 | 957 | 674 |
| 13 | Cereals | 1472 | 1582 | 1462 | 1494 |
| 14 | Beverages | 57 | 73 | 53 | 47 |
| 15 | Soft_drinks | 1374 | 1256 | 1572 | 1506 |
| 16 | Alcoholic_drinks | 375 | 475 | 458 | 135 |
| 17 | Confectionery | 54 | 64 | 62 | 41 |


I need to fix that first column...

```
rownames(x) <- x[,1]
x <- x[,1]
head(x)
```


```
[1] "Cheese"          "Carcass_meat " "Other_meat "    "Fish"
[5] "Fats_and_oils " "Sugars"
```

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names=1)
```


```
cols <- rainbow(nrow(x))
barplot(as.matrix(x), col=cols)
```



```
barplot(as.matrix(x), col=cols, beside=T)
```



```
pairs(x, col=rainbow(17), pch=16, cex=2)
```



Principal Component Analysis (PCA)

PCA can help us make sense of these types of datasets. Let's see how it works.

The main function in "base" R is called `prcomp()`. In this case we want to first take the transpose `t()` of our input `x` so the columns are the food types and the countries are the rows.

```
head(t(x))
```

| | Cheese | Carcass_meat | Other_meat | Fish | Fats_and_oils | Sugars |
|-----------|--------|--------------|------------|------|---------------|--------|
| England | 105 | 245 | 685 | 147 | 193 | 156 |
| Wales | 103 | 227 | 803 | 160 | 235 | 175 |
| Scotland | 103 | 242 | 750 | 122 | 184 | 147 |
| N.Ireland | 66 | 267 | 586 | 93 | 209 | 139 |

| | Fresh_potatoes | Fresh_Veg | Other_Veg | Processed_potatoes |
|----------|----------------|-----------|-----------|--------------------|
| England | 720 | 253 | 488 | 198 |
| Wales | 874 | 265 | 570 | 203 |
| Scotland | 566 | 171 | 418 | 220 |

| | | | | | |
|-----------|------------------|---------------|---------|-----------|-------------|
| N.Ireland | 1033 | 143 | 355 | 187 | |
| | Processed_Veg | Fresh_fruit | Cereals | Beverages | Soft_drinks |
| England | 360 | 1102 | 1472 | 57 | 1374 |
| Wales | 365 | 1137 | 1582 | 73 | 1256 |
| Scotland | 337 | 957 | 1462 | 53 | 1572 |
| N.Ireland | 334 | 674 | 1494 | 47 | 1506 |
| | Alcoholic_drinks | Confectionery | | | |
| England | 375 | 54 | | | |
| Wales | 475 | 64 | | | |
| Scotland | 458 | 62 | | | |
| N.Ireland | 135 | 41 | | | |

```
pca <- prcomp(t(x))
summary(pca)
```


Importance of components:

| | | | | |
|------------------------|----------|----------|----------|-----------|
| | PC1 | PC2 | PC3 | PC4 |
| Standard deviation | 324.1502 | 212.7478 | 73.87622 | 3.176e-14 |
| Proportion of Variance | 0.6744 | 0.2905 | 0.03503 | 0.000e+00 |
| Cumulative Proportion | 0.6744 | 0.9650 | 1.00000 | 1.000e+00 |

```
pca$x
```

| | | | | |
|-----------|------------|-------------|------------|---------------|
| | PC1 | PC2 | PC3 | PC4 |
| England | -144.99315 | -2.532999 | 105.768945 | -4.894696e-14 |
| Wales | -240.52915 | -224.646925 | -56.475555 | 5.700024e-13 |
| Scotland | -91.86934 | 286.081786 | -44.415495 | -7.460785e-13 |
| N.Ireland | 477.39164 | -58.901862 | -4.877895 | 2.321303e-13 |

```
plot( pca$x[,1], pca$x[,2], col=c("orange", "red", "blue", "darkgreen"), pch=16)
```




The “loadings” tell us how much the original variables (in our case the food) contribute to the new variables i.e. the PCs.

```
head(pca$rotation)
```

| | PC1 | PC2 | PC3 | PC4 |
|---------------|--------------|-------------|-------------|--------------|
| Cheese | -0.056955380 | 0.01601285 | 0.02394295 | -0.694538519 |
| Carcass_meat | 0.047927628 | 0.01391582 | 0.06367111 | 0.489884628 |
| Other_meat | -0.258916658 | -0.01533114 | -0.55384854 | 0.279023718 |
| Fish | -0.084414983 | -0.05075495 | 0.03906481 | -0.008483145 |
| Fats_and_oils | -0.005193623 | -0.09538866 | -0.12522257 | 0.076097502 |
| Sugars | -0.037620983 | -0.04302170 | -0.03605745 | 0.034101334 |

```
## Lets focus on PC1 as it accounts for > 90% of variance
par(mar=c(10, 3, 0.35, 0))
barplot( pca$rotation[,1], las=2 )
```



Stop at Q9